

Project 1 Writeup: Hybrid Images

Sherwyn Braganza

October 4, 2022

Summary of Contents

- The Aim of the Project
- The Approach and Algorithm
- The Results
- Project Document Questions
- Appendices

1 The Project

The aim of this project was to learn about Hybrid Images, as described by Oliva, Torralba and Schyns in [SIGGRAPH06](#). As showed in their report, the human eye tends to focus more on high frequency stuff when the object is nearby and focus on low frequency stuff when the the image is far. This can be exploited to create Hybrid Images - images that appear different when they are nearby as to when they are far. Blending filtered high-frequency content from one image on filtered low-frequency data from another image, one can create a Hybrid Image as show below (Fig 1). The following document explores this method of creating hybrid images using two methods - Spatial Domain Convolution and Fourier Domain Manipulation.




Figure 1: A hybrid image

2 The Algorithm and Approach

As mentioned in the document by Oliva et al., main process of generating hybrid images can be summarized in Equation 1 (shown below).

$$H = I_1 \otimes G_1 + I_2 \otimes (1 - G_2) \quad (1)$$

$I_1 \otimes G_1$ corresponds to the first image convoluted with a Gaussian kernel G_1 , which ends up being the low-pass filtered version of the first image. $I_2 \otimes (1 - G_2)$ corresponds to the second image being convoluted with $I_2 \otimes (1 - G_2)$ which ends up being the high-pass version of the second image. This equation represents the spatial domain method for generating the hybrid image, in the Fourier Domain, the convolution operator just turns into a normal multiplication operator. G_1 and G_2 are Gaussian Filters generated from two different values of sigma or cut-off Frequencies (Fourier Domain linguistics). The generateGaussianKernel function computes these Gaussians and its code can be found in Appendix III.

2.1 Spatial Domain Convolution

As mentioned above, the Hybrid image in the Spatial Domain is generated following Equation 1. In my implementation, I exploit the fact that image convolution with a kernel is the same as image correlation with that exact same kernel, except its flipped vertically and horizontally. The following code snippet shows that exact implementation.

```

1 def imConvolute(image: np.ndarray, kernel: np.ndarray) -> np.ndarray:
2     if kernel.shape[0] % 2 == 0 or kernel.shape[1] % 2 == 0:
3         raise Exception('Kernel with even dimensions provided.')
4
5     # flip the kernel along rows and cols
6     kernel = np.flip(np.flip(kernel, axis=1), axis=0)
7
8     # call correlation function and get the kernel correlated image as well as the impulse version of it
9     filtered = imCorrelate(image, kernel)
10
11     return filtered

```

The listing below shows my Image Correlation Function. It was designed to work almost exactly like `signal.correlate2D`. My correlate function implements correlation by linearly moving the kernel over each pixel and calculating the element wise product of the kernel element and the corresponding pixel it overlaps. The sum of these element products is the convolution result for that pixel (lines 28 - 33). Before doing that, it

pads the image with 0(s) based on the kernel size (line 15 - 20). The amount to be fairly easy to compute - its basically the half the number of rows and columns rounded down to the lower integer(line 16).

```

1 def imCorrelate(image: np.ndarray, kernel: np.ndarray) -> np.ndarray:
2     if kernel.shape[0] % 2 == 0 or kernel.shape[1] % 2 == 0:
3         raise Exception('Kernel with even dimensions provided.')
4
5     # if grayscale or colored
6     color = True if len(image.shape) > 2 else False
7
8     image = skimage.img_as_float32(image) # convert to floats in [0,1] to make computations uniform
9
10    # if grayscale, create a third dimension with only one channel
11    if not color:
12        image = image.reshape(image.shape[0], image.shape[1], 1)
13
14    # Padding section
15    pad_row, pad_col = kernel.shape[0] // 2, kernel.shape[1] // 2 # calculate pad_width for rows and cols
16    pad_params = ((pad_row, pad_row), (pad_col, pad_col), (0, 0))
17    padded_img = np.pad(image,
18                        pad_params,
19                        mode='constant',
20                        constant_values=0) # pad image along rows and cols but not channels (if it exists)
21
22    # create a container for the kernel filtered image
23    filtered = np.zeros(image.shape)
24
25    # Check if it is an rgb or grayscale img
26    channel = 3 if color else 1
27
28    for i in range(pad_row, image.shape[0] + pad_row):
29        for j in range(pad_col, image.shape[1] + pad_col):
30            for k in range(channel):
31                filtered[i - pad_row, j - pad_col, k] = np.sum(
32                    padded_img[i - pad_row:i + pad_row + 1, j - pad_col:j + pad_col + 1,
33                    k] * kernel) # convolution step
34
35    # clip images and convert them back to ubytes before returning
36
37    return skimage.img_as_ubyte(filtered.clip(0, 1)) if color \
38        else skimage.img_as_ubyte(filtered.clip(0, 1))[:, :, 0]

```

2.2 Fourier Domain Manipulation

The Algorithm followed in this method implements Equation 1 but in the Fourier. Both the original image and the hybrid image are converted to the Fourier representations of themselves using numpy's fft function (lines 20 - 22). Fc corresponds to the Forier Domain representation of th image and Fk corresponds to that of the kernel. After multiplying them both, the result is then converted back to the Spatial Domain using the ifft2 function.

```

1 def fourierDomain(image, kernel):
2     if kernel.shape[0] % 2 == 0 or kernel.shape[1] % 2 == 0:
3         raise Exception('Kernel with even dimensions provided.')
4
5     color = True if len(image.shape) > 2 else False
6
7     # if grayscale, create a third dimension with only one channel
8     if not color:
9         image = image.reshape(image.shape[0], image.shape[1], 1)
10
11    image = skimage.img_as_float32(image)
12    pad_values = (image.shape[0] - kernel.shape[0]) // 2, (image.shape[1] - kernel.shape[1]) // 2
13    padded_kernel = np.zeros(image.shape[0:2])
14    padded_kernel[pad_values[0]: pad_values[0] + kernel.shape[0], pad_values[1]: pad_values[1] + kernel.shape[1]]
15    = kernel
16
17    output_img = np.zeros(image.shape)
18
19    for i in range(image.shape[2]):
20        Fc = np.fft.fft2(np.fft.ifftshift(image[:, :, i]))
21        Fk = np.fft.fft2(np.fft.ifftshift(padded_kernel))
22        output_img[:, :, i] = np.abs(np.fft.ifftshift(np.fft.ifft2(Fc * Fk)))
23
24    return skimage.img_as_ubyte(output_img) if color else skimage.img_as_ubyte(output_img)[:, :, 0]

```

`np.fft.fft2` expects a shifted version of the image, hence the initial `ifftshift` before the conversion. It then calculates the 2 Dimensional FFT on the image and returns data with the DC part to the exterior. In order to visualize the data an `fftshift` must be performed on the data before plotting it to shift the DC component to the center. The syntax for using this function can be found [here](#). The inverse of `fft2` - `ifft2`, which converts the Fourier Data back to Spatial Representation expects shifted data, hence no `fftshift` is performed there. The result requires an extra `ifftshift` before it is plotted.

3 Results

3.1 `my_imfilter()` on different Kernels

[Figure 2](#) and [Figure 3](#) show the results of `my_imfilter()` function. It performs image convolution in the Spatial Domain using different kernels. The first image is the original, followed by the Impulse response Filter, Sobel Filter, Sharpen Filter, Emboss Filter and a Gaussian Filter. The function was also tested on an non-square kernel. The output can be found in `tests/unevenkernel.jpg`.




Figure 2: `my_imfilter` on a color image



Figure 3: `my_imfilter` on a grayscale image

3.2 Fourier Domain Tests

This section shows the results achieved by manipulating an image with a Gaussian Kernel in the Fourier Domain, at every step. Figure 4 and Figure 5 show the results from using my `fourierDomain()` function to perform convolution in the fourier Domain. Check `fourierTests()` in Appendix I for the code. The first image is the original image, the second is the Fourier Domain Spectra after calling `fft2` and `fftshift`. The third is the Fourier Domain Spectra of the image with a Gaussian Filter (low pass) imposed on it by multiplication in the Fourier Domain. The fourth is the image converted back to the spatial domain. Our results match our expectations, in which we expect to see only low frequencies.

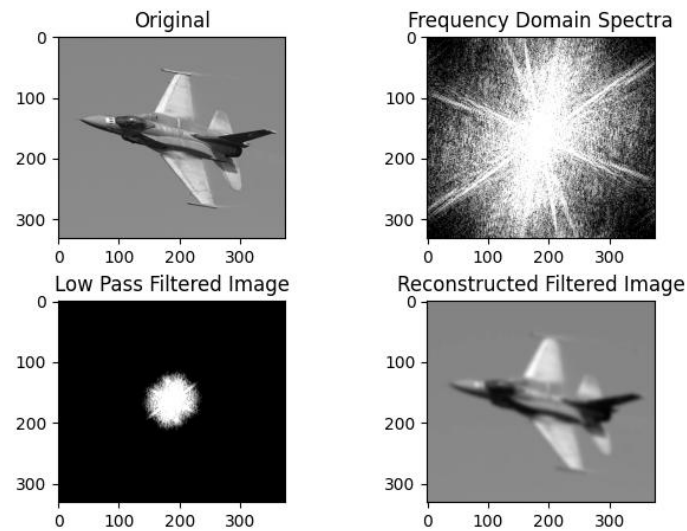


Figure 4: Fourier Domain Manipulation of plane.bmp

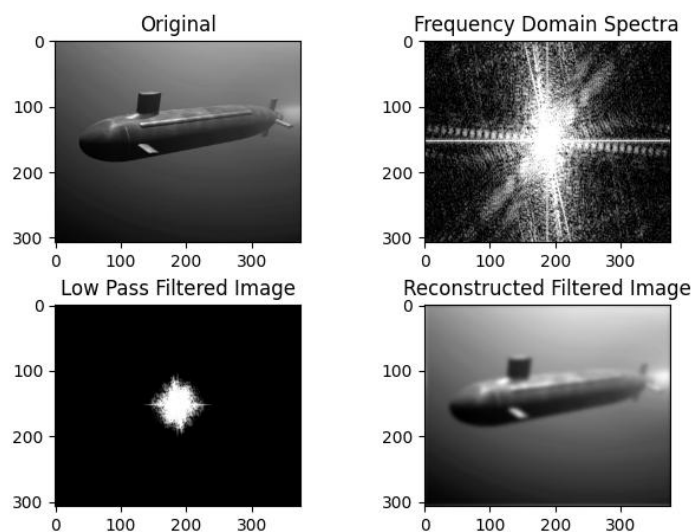


Figure 5: Fourier Domain Manipulation of submarine.bmp

3.3 Hybrid Image Generation in the Spatial Domain

This section shows the results of generating a hybrid image in the Spatial Domain. Appendix III contains the full code structure to achieve this. I have a general hybridise function and based on the value of the boolean variable `fourier`, it decides whether to use the Spatial Domain Method or the Fourier Domain Method. [Figure 6](#) shows the results achieved in generating a hybrid image in the Spatial domain using the images `cat.bmp` and `dog.bmp`.



Figure 6: Cat + Dog = Cog Hybrid Image

3.4 Hybrid Image Generation in the Fourier Domain

Like in the previous subsection, a hybrid image was generated using the same hybridise function albeit with `fourier=True`. This generated the hybrid image using the Fourier Domain Method. The image of Albert Einstein and Marilyn Monroe were melded to form the hybrid image Albert Monroe ([Figure 7](#)). From our experiments in Question 4, we saw that the Fourier Domain Method was much faster. More hybrid images were generated using this method and they can be found in Appendix IV.




Figure 7: Albert Monroe

4 Project Document Questions

Explicitly describe image convolution: the input, the transformation and the output. Why is it useful in Computer Vision?

Mathematically, the convolution operation is performed by taking the element-wise product of the first matrix with the matrix (that is flipped along its columns and rows) at each element to form a new matrix as shown below (Figure 8). The resulting matrix is usually smaller than the original and is dependent on the second matrix.

$$\begin{array}{c}
 \begin{array}{|c|c|c|c|c|c|c|c|}
 \hline 45 & 60 & 98 & 127 & 132 & 133 & 137 & 133 \\
 \hline 46 & 65 & 98 & 123 & 126 & 128 & 131 & 133 \\
 \hline 47 & 65 & 96 & 115 & 119 & 123 & 135 & 137 \\
 \hline 47 & 63 & 91 & 107 & 113 & 122 & 138 & 134 \\
 \hline 50 & 59 & 80 & 97 & 110 & 123 & 133 & 134 \\
 \hline 49 & 53 & 68 & 83 & 97 & 113 & 128 & 133 \\
 \hline 50 & 50 & 58 & 70 & 84 & 102 & 116 & 126 \\
 \hline 50 & 50 & 52 & 58 & 69 & 86 & 101 & 120 \\
 \hline
 \end{array} & * & \begin{array}{|c|c|c|}
 \hline 0.1 & 0.1 & 0.1 \\
 \hline 0.1 & 0.2 & 0.1 \\
 \hline 0.1 & 0.1 & 0.1 \\
 \hline
 \end{array} & = & \begin{array}{|c|c|c|c|c|c|c|c|}
 \hline 69 & 95 & 116 & 125 & 129 & 132 & & \\
 \hline 68 & 92 & 110 & 120 & 126 & 132 & & \\
 \hline 66 & 86 & 104 & 114 & 124 & 132 & & \\
 \hline 62 & 78 & 94 & 108 & 120 & 129 & & \\
 \hline 57 & 69 & 83 & 98 & 112 & 124 & & \\
 \hline 53 & 60 & 71 & 85 & 100 & 114 & & \\
 \hline
 \end{array} \\
 f(x,y) & & h(x,y) & & g(x,y)
 \end{array}$$

Figure 8: Convolution

Image convolution involves 2 data matrices - the image and the kernel. The image input is usually padded based on the kernel size so that the output is the same size as the original image. Convolution works in the same way as it does mathematically, each pixel value is transformed into a new value that is dependent on the pixels surrounding it. In the Fourier domain convolution turns into multiplication, which is super useful because one can visualize the kernel acting as a filter to filter out unwanted frequencies (data) of the image as seen in Figure 4 and Figure 5. This is one of the main and most useful features of convolution of images - acting as a filter to filter out unwanted frequencies (best visualized in the fourier domain).

This is super useful in Computer Vision as kernels can be designed to pick up only certain features (and frequencies) from an image, like edges and curves (Sobel Filters) whose data can then be used in Image Recognition Software or to train Deep Neural Networks (CNNs).

What is the difference between convolution and correlation? Construct a scenario which produces a different output between both operations.

Correlation and Convolution are very similar. They implement the same element-wise multiplication as explained above - the only difference is that convolution does it with the vertically and horizontally flipped version of the kernel. This results in similar, but slightly different outcomes.

In most cases where the kernel is a symmetrical matrix, convolution and correlation are the same as shown below (Table 1).

However when the kernel is not symmetrical (Sobel Filter), the kernels end up being a little different as seen in Table 2

0	-1	0
-1	5	-1
0	-1	0

0	-1	0
-1	5	-1
0	-1	0

Table 1: *Left: Correlation ; Right: Convolution*

-1	0	1
-2	0	2
-1	0	1

1	0	-1
2	0	-2
1	0	-1

Table 2: *Left: Correlation ; Right: Convolution*

What is the difference between a high pass filter and a low pass filter in how they are constructed, and what they do to the image? Please provide example kernels and output images.

High Pass Filters allow only high frequency data through while blocking low frequency data; a Low Pass Filter is the converse of that.

A common low pass filter is a Gaussian or Normal Distribution Filter with a mean of 0. The Fourier Spectrum of this Image turns out to be a circle with 0s on the outside and positive values in the inside, which when multiplied with another image creates a low pass filtered version of that image as we know that all lower frequency content lies close to the center. (shown below and in [Figure 4](#)).

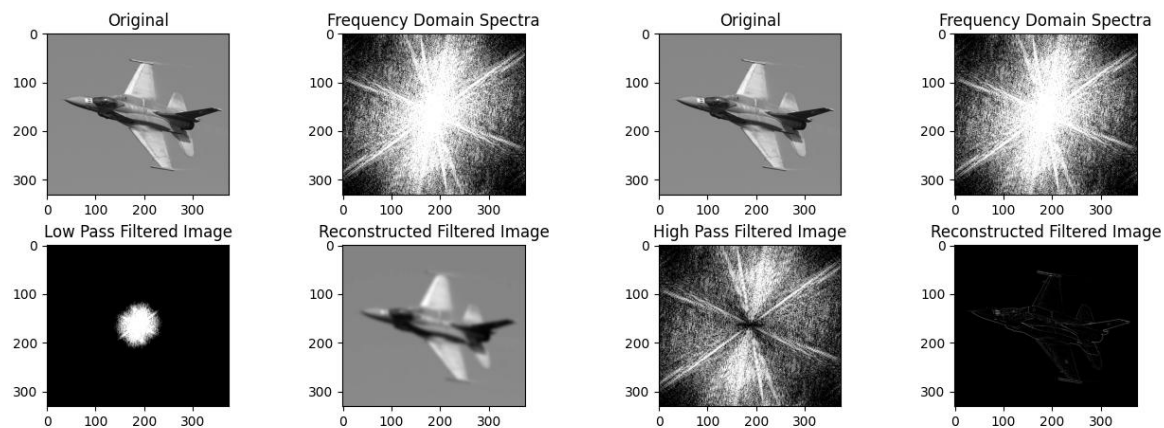


Figure 9: Low Pass and High Pass Filtered plane.bmp

Subtracting this newly obtained low frequency image from the original image in the Fourier Domain, results in a high frequency version of that image. From this we can infer that a high pass filter is just $1 - G$ where G is a low-pass Gaussian Filter. The figure above shows this in play ([Figure 9](#)).

How does computation time vary with filter sizes from 3×3 to 15×15 (for all odd and square sizes), and with image sizes from 0.25 MPix to 8 MPix (choose your own intervals - you can use the image in project 0 if you'd like)? Measure both using `imfilter` to produce a matrix of values. Use the `imresize` function to vary the size of an image. Use an appropriate charting function to plot your matrix of results, such as `scatter3` or `surf`. Do the results match your expectation given the number of multiply and add operations in convolution?

For the most part the results matched my expectations - the computation time should increase non linearly with the increase in image size and kernel size. My results are show below in [Figure 10](#).




Figure 10: Image Size and Kernel Size Computation Metrics

Apart from plotting Convolution computation times, I also plotted Fourier Domain Manipulation Convolution times. Fourier Domain Manipulation performed almost as good Convolution for smaller kernels and smaller images but outshone convolution with when the image sizes and kernel sizes got bigger. However what was surprising was that Convolution performed better than Fourier Domain Manipulation when sizes of images were small but kernels were big.

5 Appendix I

```

1 import os
2 import numpy as np
3 import skimage
4 from skimage import io
5 import matplotlib.pyplot as plt
6
7 IMPULSE = np.asarray([[0,0,0],
8                       [0,1,0],
9                       [0,0,0]])
10 SOBEL = np.asarray([[ -1, 0, 1],
11                    [-2, 0, 2],
12                    [-1, 0, 1]])
13 UNEVEN_SOBEL = np.asarray([[ -2, -1, 0, 1, 2],
14                           [-4, -2, 0, 2, 4],
15                           [-2, -1, 0, 1, 2]])
16 SHARPEN = np.array([[0, -1, 0],
17                    [-1, 5, -1],
18                    [0, -1, 0]])
19 EMOSS = np.array([[ -2, -1, 0],
20                  [-1, 1, 1],
21                  [0, 1, 2]])
22
23
24 def my_imfilter(image: np.ndarray, kernel: np.ndarray) -> (np.ndarray, np.ndarray):
25     """
26     Wrapper Function for imConvolute. Created to meet assignment
27     document specifications
28     """
29     return imConvolute(image, IMPULSE), imConvolute(image, kernel)
30
31
32 def imConvolute(image: np.ndarray, kernel: np.ndarray) -> np.ndarray:
33     """
34     Image Convolution Filter
35
36     Author: Sherwyn Braganza
37     Sept 28, 2022 - Initial Creation
38     Sept 28, 2022 - Added more descriptive comments.
39     Sept 28, 2022 - Finished off most of convolve
40     Sept 28, 2022 - Compared results with signal.convolve2D
41     Sept 30, 2022 - Changes made to return only filtered image
42
43     Function that implements convolution through correlation. Designed to work
44     like signal.convolve2D from the scipy library. The kernel used for convolution
45     is presented by the user. This function is designed to handle both rgb and grayscale images.
46
47     It exploits the fact that convolution is basically correlation using a flipped kernel.
48     It therefore flips the kernel and passes it to the function that gives the correlational
49     version of an image with the flipped kernel.
50
51     The function returns a np.ndarray, corresponding to the image correlated with the kernel provided.
52
53     :param image: Numpy matrix containing image data
54     :param kernel: Numpy matrix containing data for an
55                   odd dimension kernel
56
57     :return filter: Kernel Convolved Version of the Image
58     """
59     if kernel.shape[0] % 2 == 0 or kernel.shape[1] % 2 == 0:
60         raise Exception('Kernel with even dimensions provided.')
61
62     # flip the kernel along rows and cols
63     kernel = np.flip(np.flip(kernel, axis=1), axis=0)
64
65     # call correlation function and get the kernel correlated image as well as the impulse version of it
66     filtered = imCorrelate(image, kernel)
67
68     return filtered
69
70
71 def imCorrelate(image: np.ndarray, kernel: np.ndarray) -> np.ndarray:
72     """
73     Image Correlation Filter
74
75     Author: Sherwyn Braganza
76     Sept 28, 2022 - Initial Creation
77     Sept 28, 2022 - Added more descriptive comments.
78     Sept 28, 2022 - Finished off most of correlate
79     Sept 28, 2022 - Compared results with signal.correlate2D
80     Sept 30, 2022 - Changes made to return only filtered image
81     Sept 30, 2022 - Increased compatibility with grayscales
82
83     Function that implements correlation image processing. Designed to work
84     like signal.correlate2D from the scipy library. The kernel used for correlation
85     is presented by the user.
86
87     This function convolves the kernel with the image using matrix convolution. It
88     initially converts the image to a float and then pads it with 0s along the borders

```

```

89     according to the shape of the convolutional kernel. It finally clips the image pixel values to [0,1]
90     before converting it back to ubyte format and returning it.
91
92     The function returns a np.ndarray, corresponding to the image correlated with the kernel provided.
93
94     TODO – Try to use logical indexing insted of loops
95
96     :param image: Numpy matrix containing image data
97     :param kernel: Numpy matrix containing data for an odd dimension kernel
98
99     :return filter: Kernel Correlated Version of the Image
100 """
101 if kernel.shape[0] % 2 == 0 or kernel.shape[1] % 2 == 0:
102     raise Exception('Kernel with even dimensions provided.')
103
104 # if grayscale or colored
105 color = True if len(image.shape) > 2 else False
106
107 image = skimage.img_as_float32(image) # convert to floats in [0,1] to make computations uniform
108
109 # if grayscale, create a third dimension with only one channel
110 if not color:
111     image = image.reshape(image.shape[0], image.shape[1], 1)
112
113 # Padding section
114 pad_row, pad_col = kernel.shape[0] // 2, kernel.shape[1] // 2 # calculate pad-width for rows and cols
115 pad_params = ((pad_row, pad_row), (pad_col, pad_col), (0, 0))
116 padded_img = np.pad(image,
117                     pad_params,
118                     mode='constant',
119                     constant_values=0) # pad image along rows and cols but not channels (if it exists)
120
121 # create a container for the kernel filtered image
122 filtered = np.zeros(image.shape)
123
124 # Check if it is an rgb or grayscale img
125 channel = 3 if color else 1
126
127 for i in range(pad_row, image.shape[0] + pad_row):
128     for j in range(pad_col, image.shape[1] + pad_col):
129         for k in range(channel):
130             filtered[i - pad_row, j - pad_col, k] = np.sum(
131                 padded_img[i - pad_row:i + pad_row + 1, j - pad_col:j + pad_col + 1,
132                 k] * kernel) # convolution step
133
134 # clip images and convert them back to bytes before returning
135
136 return skimage.img_as_ubyte(filtered.clip(0, 1)) if color \
137     else skimage.img_as_ubyte(filtered.clip(0, 1))[:, :, 0]
138
139
140 def hybridise(image1: np.ndarray, image2: np.ndarray, sigma1: float, sigma2: float, fourier: bool):
141     """
142     Hybrid Image Generator
143
144     Generates a hybrid image according to the process described by Oliva, Torralba and Schyns
145     in Siggraph (2006) (http://olivalab.mit.edu/hybrid/Talk-Hybrid-Siggraph06.pdf).
146
147     The basis of this process is to take the high pass version of one image and
148     superimpose it on the low pass version of the other. These images have to be normalized
149     and centered for best results. This function implements the same process
150     using 2 approaches – Spatial Convolution and Fourier Domain Multiplication (Hadamard Product)
151
152     The function generates 2 different Gaussian Filters based on the sigmas provided and
153     uses them to get a low pass and high pass filtered images
154
155     The generated hybrid image is rescaled 4 times and stacked horizontally to get the
156     final image that is then returned.
157
158     Author: Sherwyn Braganza
159     29 Sept, 2022 – Implemented a rudimentary hybrid image generator
160                   that uses Spatial Convolution
161     30 Sept, 2022 – Implemented hybrid image rescaling and stacking
162     30 Sept, 2022 – Expanded it to use Spatial Convolution and
163                   Fourier Domain Hadamard (fourier domain still need to be coded up)
164
165     :param image1: The low pass intended image
166     :param image2: The high pass intended image
167     :param sigma1: The sigma value corresponding to the low pass image
168     :param sigma2: The sigma value corresponding to the high pass image
169     :param fourier: False if you want to use Spatial Convolution, True
170                   if you want to use Fourier Domain Processing
171
172     :return: The stacked hybrid image
173     """
174     hybrid = np.asarray([])
175     gaussian_low = generateGaussianKernel(sigma1)
176     gaussian_high = generateGaussianKernel(sigma2)
177     lowpass_image = lowpass_image2 = None
178
179     if not fourier:

```

```

181     lowpass_image = imConvolute(image1, gaussian_low)
182     lowpass_image2 = imConvolute(image2, gaussian_high)
183     else:
184         lowpass_image = fourierDomain(image1, gaussian_low)
185         lowpass_image2 = fourierDomain(image2, gaussian_high)
186
187     hybrid = skimage.img_as_ubyte((
188         skimage.img_as_float32(lowpass_image) +
189         skimage.img_as_float32(image2) -
190         skimage.img_as_float32(lowpass_image2)
191     ).clip(0, 1))
192
193     # if grayscale, create a third dimension with only one channel
194     if len(hybrid.shape) < 3:
195         hybrid = hybrid.reshape(hybrid.shape[0], hybrid.shape[1], 1)
196
197     #####
198     # Image stacking and padding section
199     #####
200     image_stack = [hybrid]
201
202     # create scale down versions of the original
203     for i in range(0, 4):
204         image_stack.append(skimage.img_as_ubyte(
205             skimage.transform.rescale(image_stack[i], 0.5, anti_aliasing=True, channel_axis=2)
206         ))
207
208     # padding the rescaled images along axis 0 to make them the same size vertically
209     for i in range(1, 5):
210         image_stack[i] = np.pad(
211             image_stack[i],
212             ((image_stack[0].shape[0] - image_stack[i].shape[0], 0),
213              (0, 0),
214              (0, 0)),
215             mode='constant',
216             constant_values=255
217         )
218
219     # padding along axis 1
220     for i in range(1, 5):
221         image_stack[i] = np.pad(
222             image_stack[i],
223             ((0, 0),
224              (5, 0),
225              (0, 0)),
226             mode='constant',
227             constant_values=255
228         )
229
230     return np.hstack(image_stack) if len(image1.shape) > 2 else np.hstack(image_stack)[:,:,:0]
231
232
233 def fourierDomain(image, kernel):
234     """
235     Fourier Domain Image Convolution
236
237     Author: Sherywn Braganza
238
239     Implements image convolution by shifting the image and kernel into the Frequency Domain, computing
240     the Hadamard Product of them both and then converting them back to the spacial domain using the
241     inverse fourier transform.
242
243     :param image: The image to be convolved
244     :param kernel: The kernel to be convolved
245     :return: The convolved image and kernel
246     """
247     if kernel.shape[0] % 2 == 0 or kernel.shape[1] % 2 == 0:
248         raise Exception('Kernel with even dimensions provided.')
249
250     color = True if len(image.shape) > 2 else False
251
252     # if grayscale, create a third dimension with only one channel
253     if not color:
254         image = image.reshape(image.shape[0], image.shape[1], 1)
255
256     image = skimage.img_as_float32(image)
257     pad_values = (image.shape[0] - kernel.shape[0]) // 2, (image.shape[1] - kernel.shape[1]) // 2
258     padded_kernel = np.zeros(image.shape[0:2])
259     padded_kernel[pad_values[0]: pad_values[0] + kernel.shape[0], pad_values[1]: pad_values[1] + kernel.shape[1]] = kernel
260
261     output_img = np.zeros(image.shape)
262
263     for i in range(image.shape[2]):
264         Fc = np.fft.fft2(np.fft.ifftshift(image[:, :, i]))
265         Fk = np.fft.fft2(np.fft.ifftshift(padded_kernel))
266         output_img[:, :, i] = np.abs(np.fft.ifftshift(np.fft.ifft2(Fc * Fk))).clip(-1, 1)
267
268     return skimage.img_as_ubyte(output_img) if color else skimage.img_as_ubyte(output_img)[:,:,:0]
269
270
271 def generateGaussianKernel(sigma: float) -> np.ndarray:

```

```

273 """
274     Generates a 2-D Gaussian Kernel
275
276     Author: Sherwyn Braganza
277     Sept 29, 2020 - Added function and base code for generating it
278     Sept 29, 2020 - Implemented a weighted mean based kernel generator
279     Sept 30, 2020 - Changed implementation to generate a true gaussian
280                     based kernel
281
282     Generates a 1D Gaussian distribution from the Gaussian equation
283     using the value of sigma. Matrix multiplies the transpose of
284     the 1D Gaussian with itself to form a 2D square Gaussian kernel
285
286     :param sigma: The standard deviation of the gaussian
287     :return: kernel: The 2D Gaussian kernel generated in float format
288 """
289 size = int(8 * sigma + 1)
290 # enforce an odd sized kernel
291 if not size % 2:
292     size = size + 1
293
294 center = size // 2
295 kernel = np.zeros(size)
296
297 # Generate Gaussian blur.
298 for x in range(size):
299     diff = (x - center) ** 2
300     kernel[x] = np.exp(-diff / (2 * sigma ** 2))
301
302 kernel = np.asarray(kernel.reshape(-1, 1).T * kernel.reshape(-1, 1))
303 kernel = kernel / np.sum(kernel)
304
305 return kernel
306
307
308 def testConvolutionColor():
309     """
310         Script to test out Colored Image Convolution
311
312         Author: Sherwyn Braganza
313
314         :param: NONE
315         :return: NONE
316     """
317
318     img1 = io.imread('data/dog.bmp', as_gray=False)
319     GAUSSIAN = generateGaussianKernel(3)
320     img_impulse = imConvolute(img1, IMPULSE)
321     img_sobel = imConvolute(img1, SOBEL)
322     img_sharpen = imConvolute(img1, SHARPEN)
323     img_emboss = imConvolute(img1, EMBOSS)
324     img_gaussian = imConvolute(img1, GAUSSIAN)
325
326     joined = np.hstack((img1, img_impulse, img_sobel, img_sharpen, img_emboss, img_gaussian))
327     fig, axs = plt.subplots()
328     axs.set_title('Original → Impulse → Sobel → Sharpen → Emboss → Gaussian')
329     axs.imshow(joined)
330     fig.savefig('tests/my-filter-test-colored.jpg')
331
332
333 def testConvolutionGray():
334     """
335         Script to test out Colored Grayscale Convolution
336
337         Author: Sherwyn Braganza
338
339         :param: NONE
340         :return: NONE
341     """
342
343     img1 = skimage.img_as_ubyte(io.imread('data/marylin.bmp', as_gray=True))
344     GAUSSIAN = generateGaussianKernel(3)
345     img_impulse = imConvolute(img1, IMPULSE)
346     img_sobel = imConvolute(img1, SOBEL)
347     img_sharpen = imConvolute(img1, SHARPEN)
348     img_emboss = imConvolute(img1, EMBOSS)
349     img_gaussian = imConvolute(img1, GAUSSIAN)
350
351     joined = np.hstack((img1, img_impulse, img_sobel, img_sharpen, img_emboss, img_gaussian))
352     fig, axs = plt.subplots()
353     axs.set_title('Original → Impulse → Sobel → Sharpen → Emboss → Gaussian')
354     axs.imshow(joined, cmap='gray')
355     fig.savefig('tests/my-filter-test-gray.jpg')
356
357
358 def testFFT(image_name):
359     """
360         Test image processing in the Fourier Space
361
362         Converts the image to its fourier space representation, applies a gaussian
363         filter to it and then converts it back to the spatial domain. Saves images
364         at each of these steps.

```

```

365         :param image_name: Name of the image to test
366         :return: None
367     """
368     image = io.imread(image_name, as_gray=True)
369     fig, axs = plt.subplots(2, 2)
370     fig.tight_layout(h_pad=2)
371     plt.subplots_adjust(top=0.9)
372     padded_gaussian = np.zeros(image.shape)
373     gaussian = generateGaussianKernel(3)
374     pad_params = (image.shape[0] - gaussian.shape[0]) // 2, (image.shape[1] - gaussian.shape[1]) // 2
375     padded_gaussian[pad_params[0]:pad_params[0] + gaussian.shape[0],
376                    pad_params[1]:pad_params[1] + gaussian.shape[1]] = gaussian
377
378     axs[0, 0].imshow(image, cmap='gray')
379     axs[0, 0].set_title('Original')
380
381     img_fft = np.fft.fft2(np.fft.ifftshift(image))
382     axs[0, 1].imshow(skimage.img_as_ubyte(np.real(np.log10(np.fft.fftshift(img_fft))).clip(-1, 1)), cmap='gray')
383     axs[0, 1].set_title('Frequency Domain Spectra')
384
385     filtered_fft = img_fft * np.fft.fft2(np.fft.ifftshift(padded_gaussian))
386     axs[1, 0].imshow(skimage.img_as_ubyte(np.real(np.log10(np.fft.fftshift(filtered_fft))).clip(-1, 1)), cmap='gray')
387     axs[1, 0].set_title('Low Pass Filtered Image')
388
389     reconstructed = np.real(np.fft.ifftshift(np.fft.ifft2(filtered_fft))).clip(-1, 1)
390     axs[1, 1].imshow(skimage.img_as_ubyte(reconstructed), cmap='gray')
391     axs[1, 1].set_title('Reconstructed Filtered Image')
392
393     fig.savefig('tests/fft-transformed-'+image_name[5:-3]+'jpg')
394
395
396 def testUnevenKernel():
397     # Uneven Kernel Test
398     img1 = skimage.img_as_ubyte(io.imread('data/marilyn.bmp', as_gray=True))
399     convolved = imConvolute(img1, UNEVEN.SOBEL)
400     fig, axs = plt.subplots()
401     axs.set_title('Uneven Kernel')
402     axs.imshow(convolved, cmap='gray')
403     fig.savefig('tests/unevenkernel.jpg')
404
405
406 if __name__ == '__main__':
407     # check if directories exist
408     if not os.path.exists('data'):
409         os.makedirs('data')
410     if not os.path.exists('tests'):
411         os.makedirs('tests')
412
413     testConvolutionColor()
414     testConvolutionGray()
415     testFFT("data/submarine.bmp")
416     testFFT("data/plane.bmp")
417     testUnevenKernel()
418

```

6 Appendix II

```

1 def generateGaussianKernel(sigma: float) -> np.ndarray:
2     """
3     Generates a 2-D Gaussian Kernel
4
5     Author: Sherwyn Braganza
6     Sept 29, 2020 - Added function and base code for generating it
7     Sept 29, 2020 - Implemented a weighted mean based kernel generator
8     Sept 30, 2020 - Changed implementation to generate a true gaussian
9                     based kernel
10
11     Generates a 1D Gaussian distribution from the Gaussian equation
12     using the value of sigma. Matrix multiplies the transpose of
13     the 1D Gaussian with itself to form a 2D square Gaussian kernel
14
15     :param sigma: The standard deviation of the gaussian
16     :return: kernel: The 2D Gaussian kernel generated in float format
17     """
18     size = int(8 * sigma + 1)
19     # enforce an odd sized kernel
20     if not size % 2:
21         size = size + 1
22
23     center = size // 2
24     kernel = np.zeros(size)
25
26     # Generate Gaussian blur.
27     for x in range(size):
28         diff = (x - center) ** 2
29         kernel[x] = np.exp(-diff / (2 * sigma ** 2))
30
31     kernel = np.asarray(kernel.reshape(-1, 1).T * kernel.reshape(-1, 1))
32     kernel = kernel / np.sum(kernel)
33
34     return kernel

```


7 Appendix III

```

1 def hybridise(image1: np.ndarray, image2: np.ndarray, sigma1: float, sigma2: float, fourier: bool):
2     hybrid = np.asarray([])
3     gaussian_low = generateGaussianKernel(sigma1)
4     gaussian_high = generateGaussianKernel(sigma2)
5     lowpass_image = lowpass_image2 = None
6
7
8     if not fourier:
9         lowpass_image = imConvolute(image1, gaussian_low)
10        lowpass_image2 = imConvolute(image2, gaussian_high)
11    else:
12        lowpass_image = fourierDomain(image1, gaussian_low)
13        lowpass_image2 = fourierDomain(image2, gaussian_high)
14
15    hybrid = skimage.img_as_ubyte((
16        skimage.img_as_float32(lowpass_image) +
17        skimage.img_as_float32(image2) -
18        skimage.img_as_float32(lowpass_image2)
19    ).clip(0, 1))
20
21    # if grayscale, create a third dimension with only one channel
22    if len(hybrid.shape) < 3:
23        hybrid = hybrid.reshape(hybrid.shape[0], hybrid.shape[1], 1)
24
25    #####
26    # Image stacking and padding section
27    #####
28    image_stack = [hybrid]
29
30    # create scale down versions of the original
31    for i in range(0, 4):
32        image_stack.append(skimage.img_as_ubyte(
33            skimage.transform.rescale(image_stack[i], 0.5, anti_aliasing=True, channel_axis=2)
34        ))
35
36    # padding the rescaled images along axis 0 to make them the same size vertically
37    for i in range(1, 5):
38        image_stack[i] = np.pad(
39            image_stack[i],
40            ((image_stack[0].shape[0] - image_stack[i].shape[0], 0),
41             (0, 0),
42             (0, 0)),
43            mode='constant',
44            constant_values=255
45        )
46
47    # padding along axis 1
48    for i in range(1, 5):
49        image_stack[i] = np.pad(
50            image_stack[i],
51            ((0, 0),
52             (5, 0),
53             (0, 0)),
54            mode='constant',
55            constant_values=255
56        )
57
58    return np.hstack(image_stack) if len(image1.shape) > 2 else np.hstack(image_stack)[:,: ,0]

```

8 Appendix IV



Figure 11: Bicycle + Motorcycle = Bimotorcycle




Figure 12: Plane + Bird = Plird




Figure 13: Fish + Submarine = Fishmarine